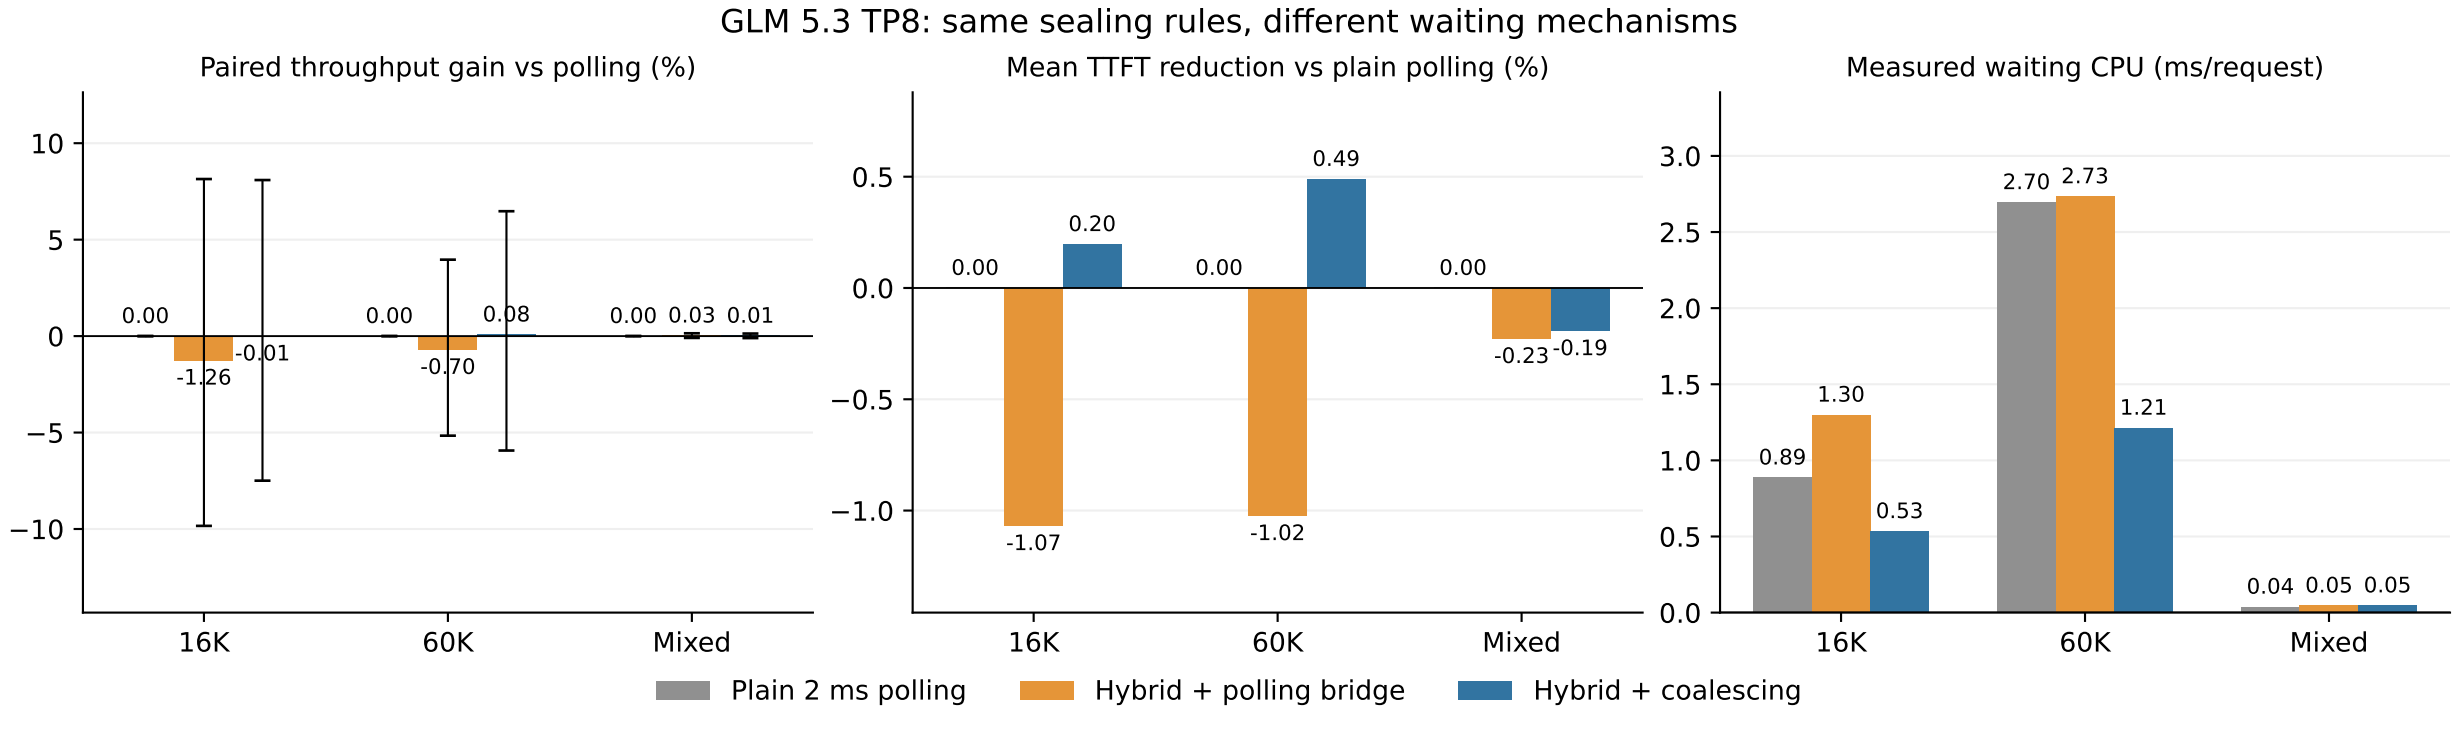
GLM 5.3 TP8: same sealing rules, different waiting mechanisms
Paired throughput gain vs polling (%)
Mean TTFT reduction vs plain polling (%)
Measured waiting CPU (ms/request)
0.00
-1.26
-0.01
0.00
-0.70
0.08
0.00
0.03
0.01
0.00
0.20
0.00
0.49
0.00
-0.23
-0.19
-1.07
-1.02
0.89
1.30
0.53
2.70
2.73
1.21
0.04
0.05
0.05
16K
60K
Mixed
Plain 2 ms polling
Hybrid + polling bridge
Hybrid + coalescing