Final strict confirmation · eight B300 GPUs · no model compute

4K prefix · 4,096 tokens
195 MB per load per GPU

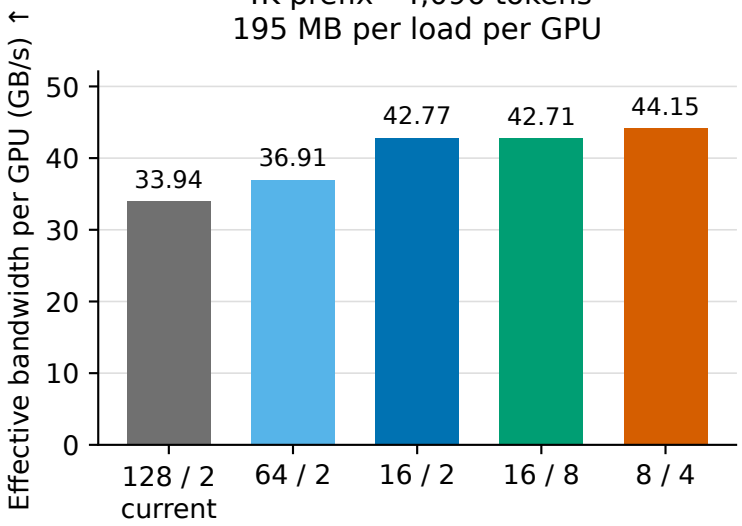

16K prefix · 15,552 tokens
742 MB per load per GPU

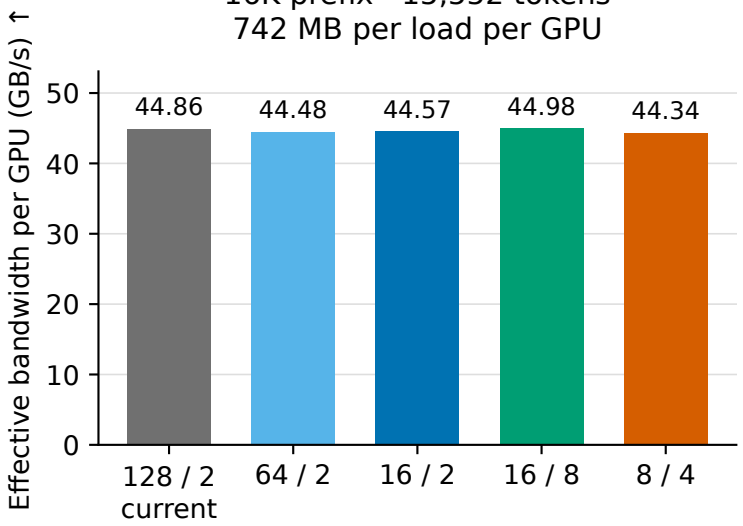

60K prefix · 58,368 tokens
2.784 GB per load per GPU

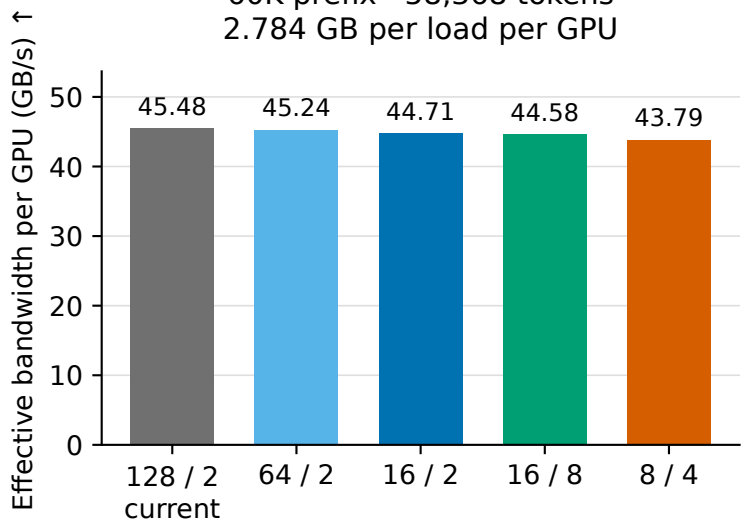

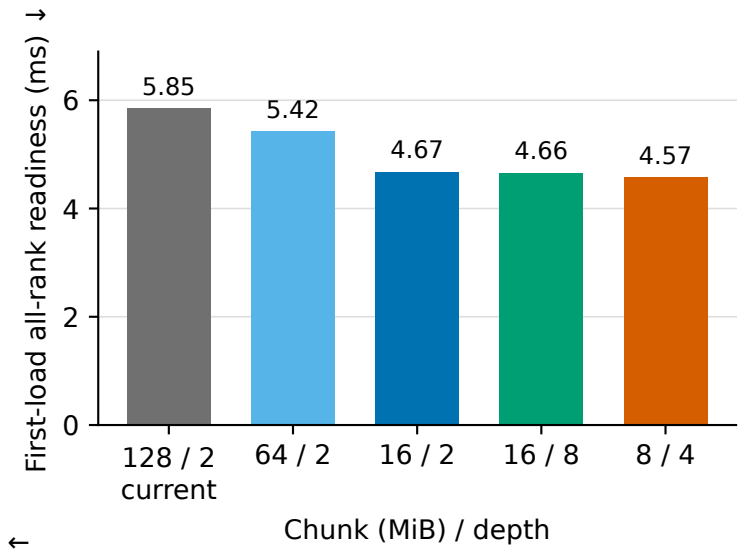

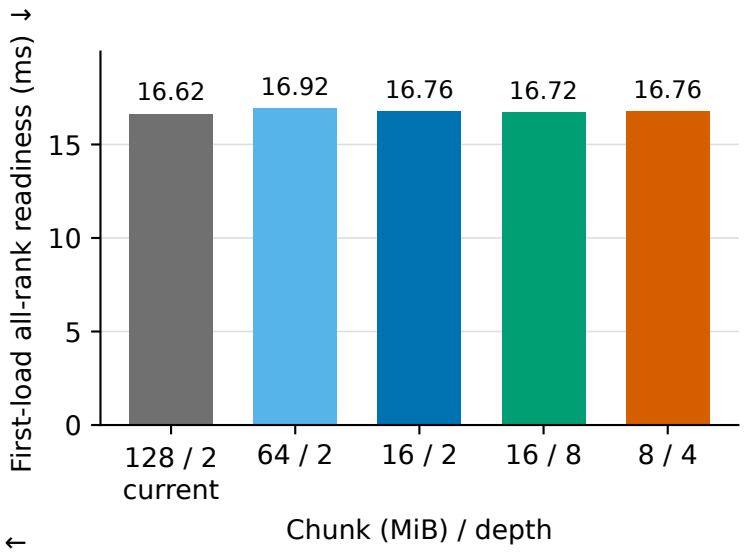

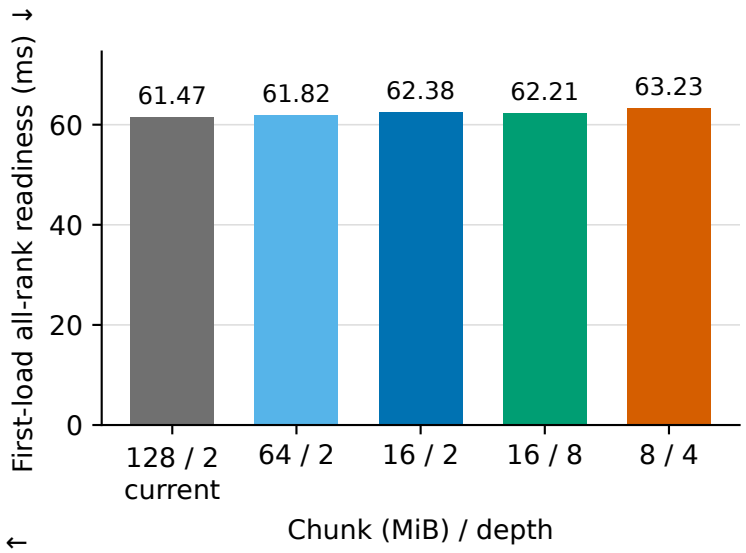

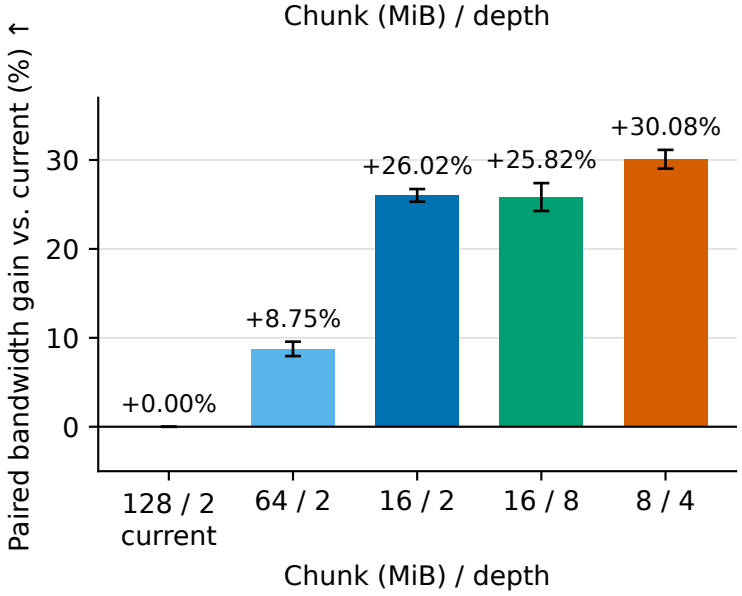

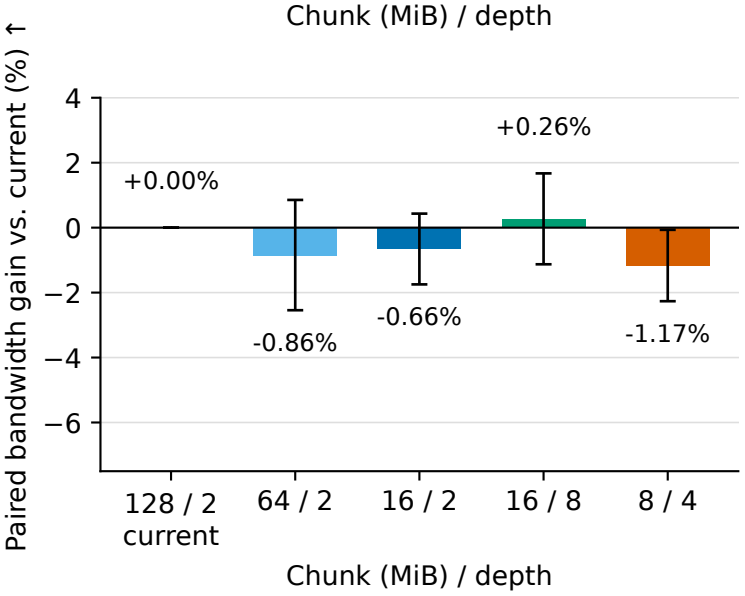

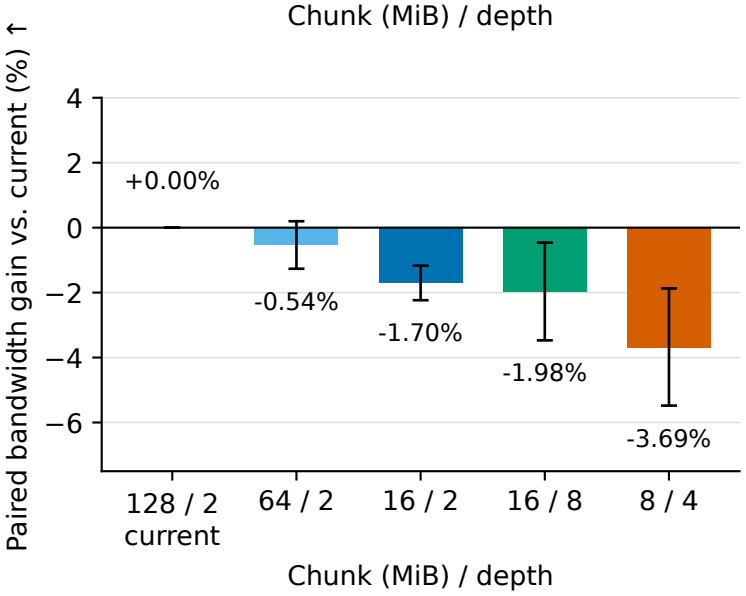

Five preselected configurations × 3 sizes × 8 repeats; validation waits for all ranks to finish timed copies.
Error bars: nominal paired 95% intervals, unadjusted for multiple comparisons. Transfer bandwidth is not serving throughput.