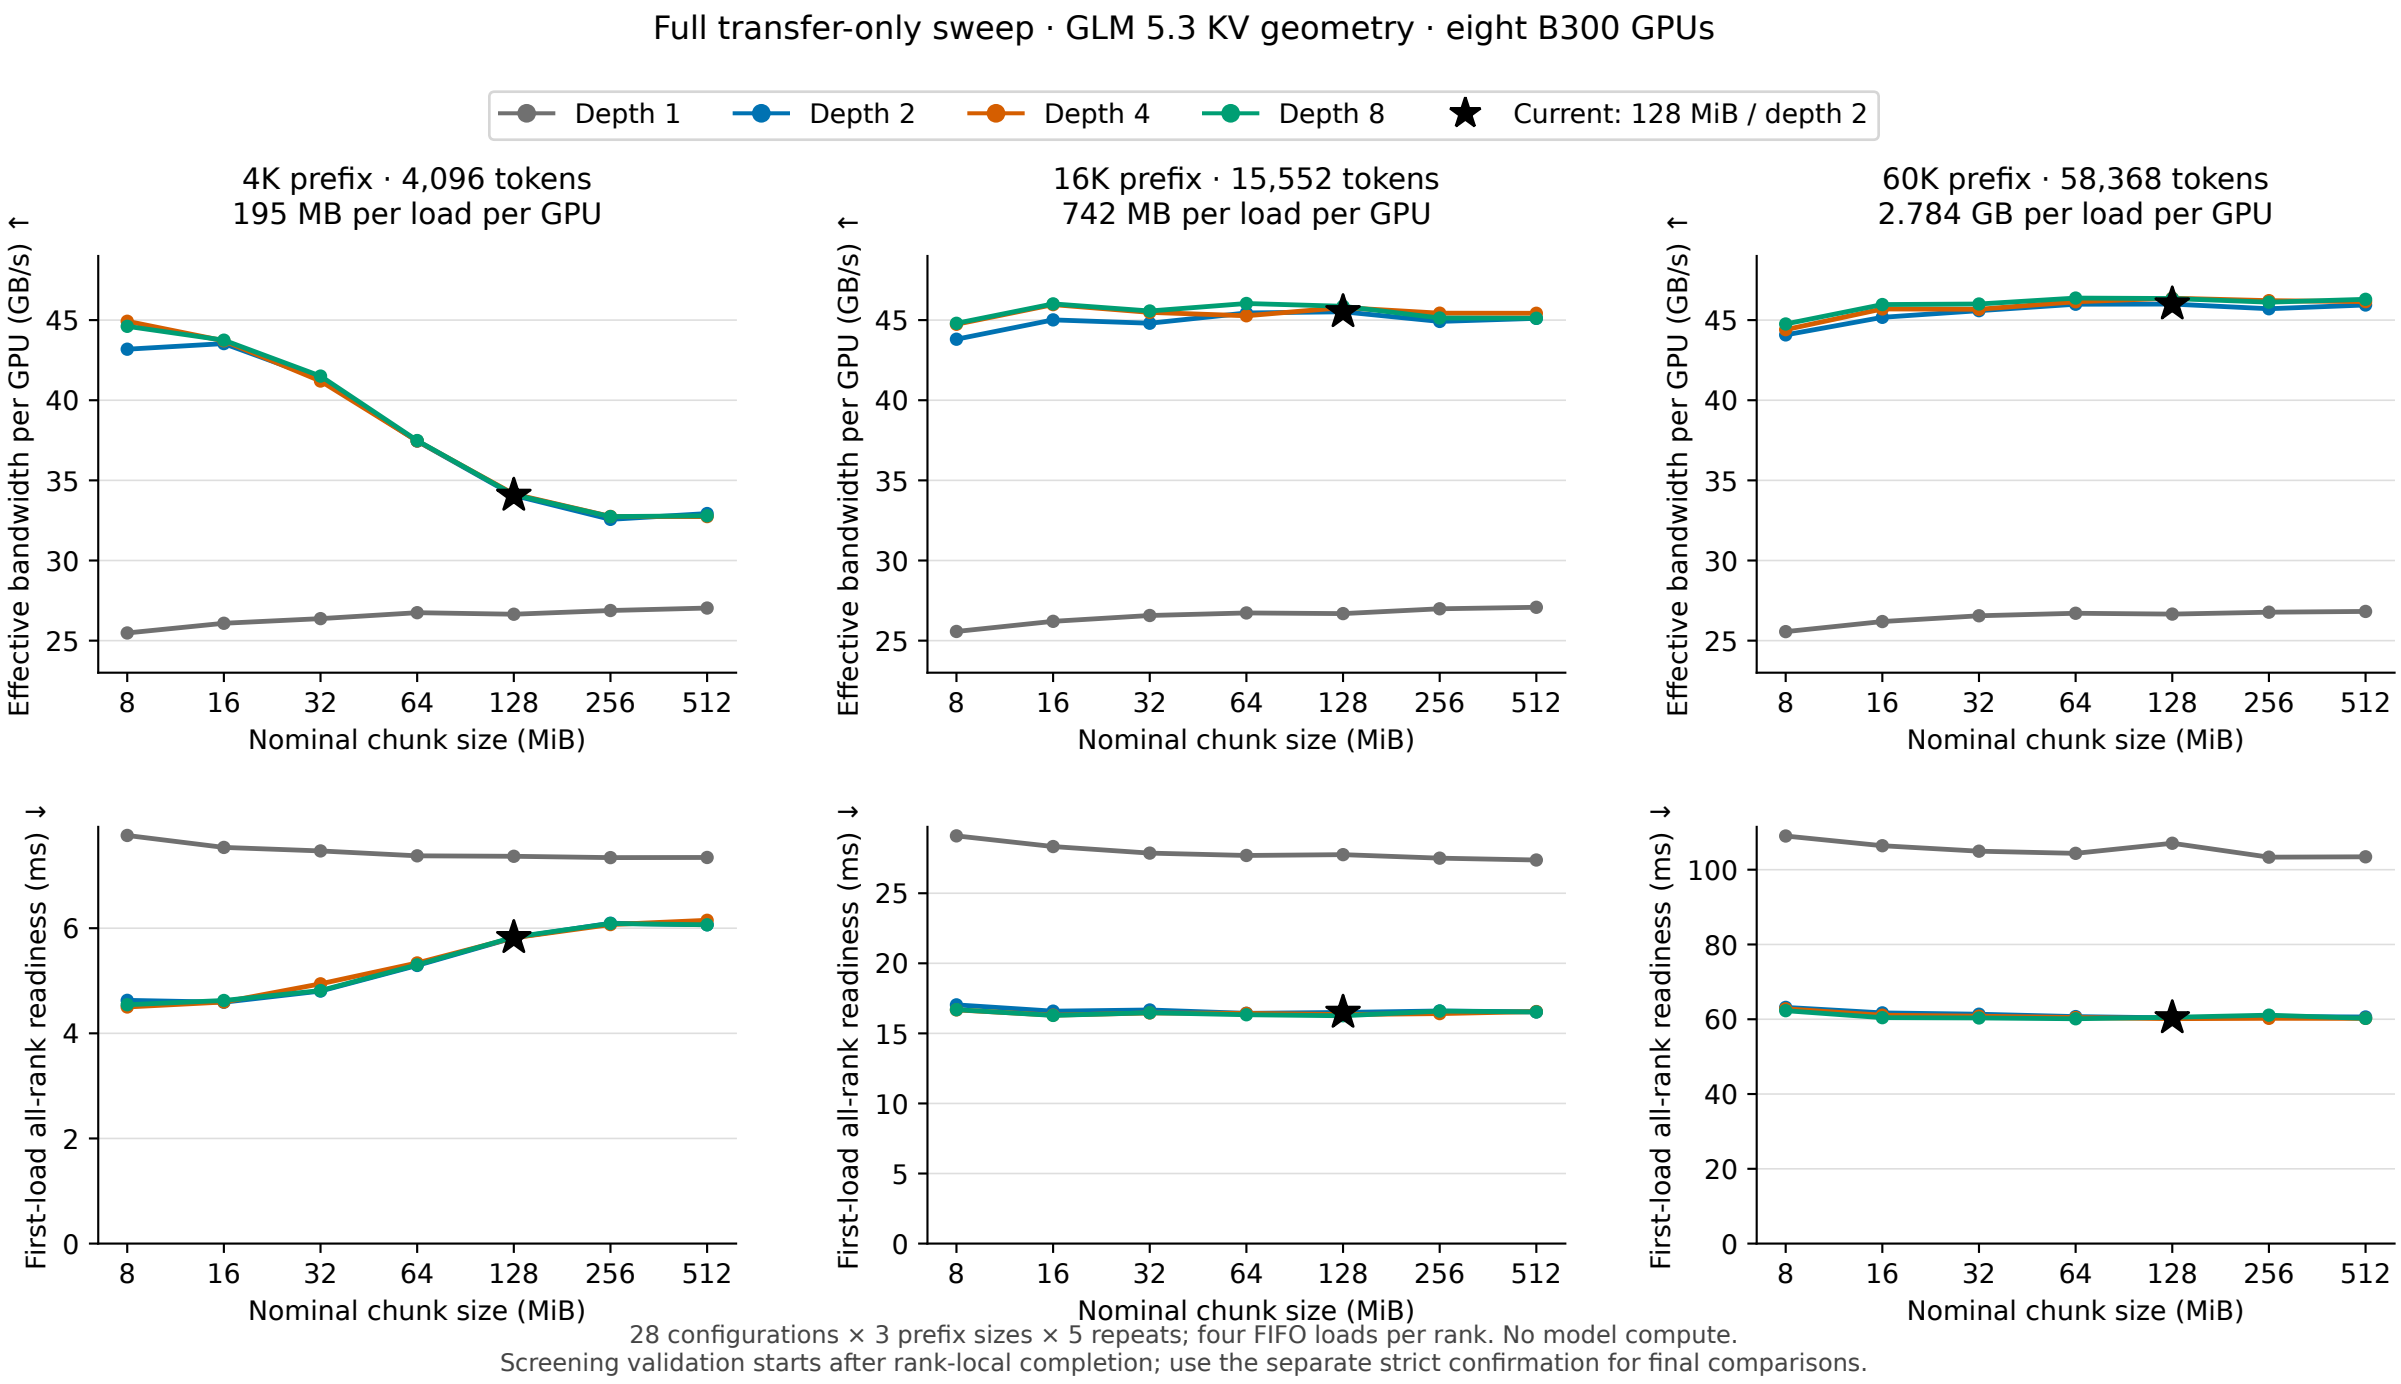
Full transfer-only sweep · GLM 5.3 KV geometry · eight B300 GPUs
Depth 1
Depth 2
Depth 4
Depth 8
Current: 128 MiB / depth 2
4K prefix · 4,096 tokens
195 MB per load per GPU
16K prefix · 15,552 tokens
742 MB per load per GPU
60K prefix · 58,368 tokens
2.784 GB per load per GPU
Effective bandwidth per GPU (GB/s) ↑
First-load all-rank readiness (ms) ↓
Nominal chunk size (MiB)
8
16
32
64
128
256
512
28 configurations × 3 prefix sizes × 5 repeats; four FIFO loads per rank. No model compute.
Screening validation starts after rank-local completion; use the separate strict confirmation for final comparisons.