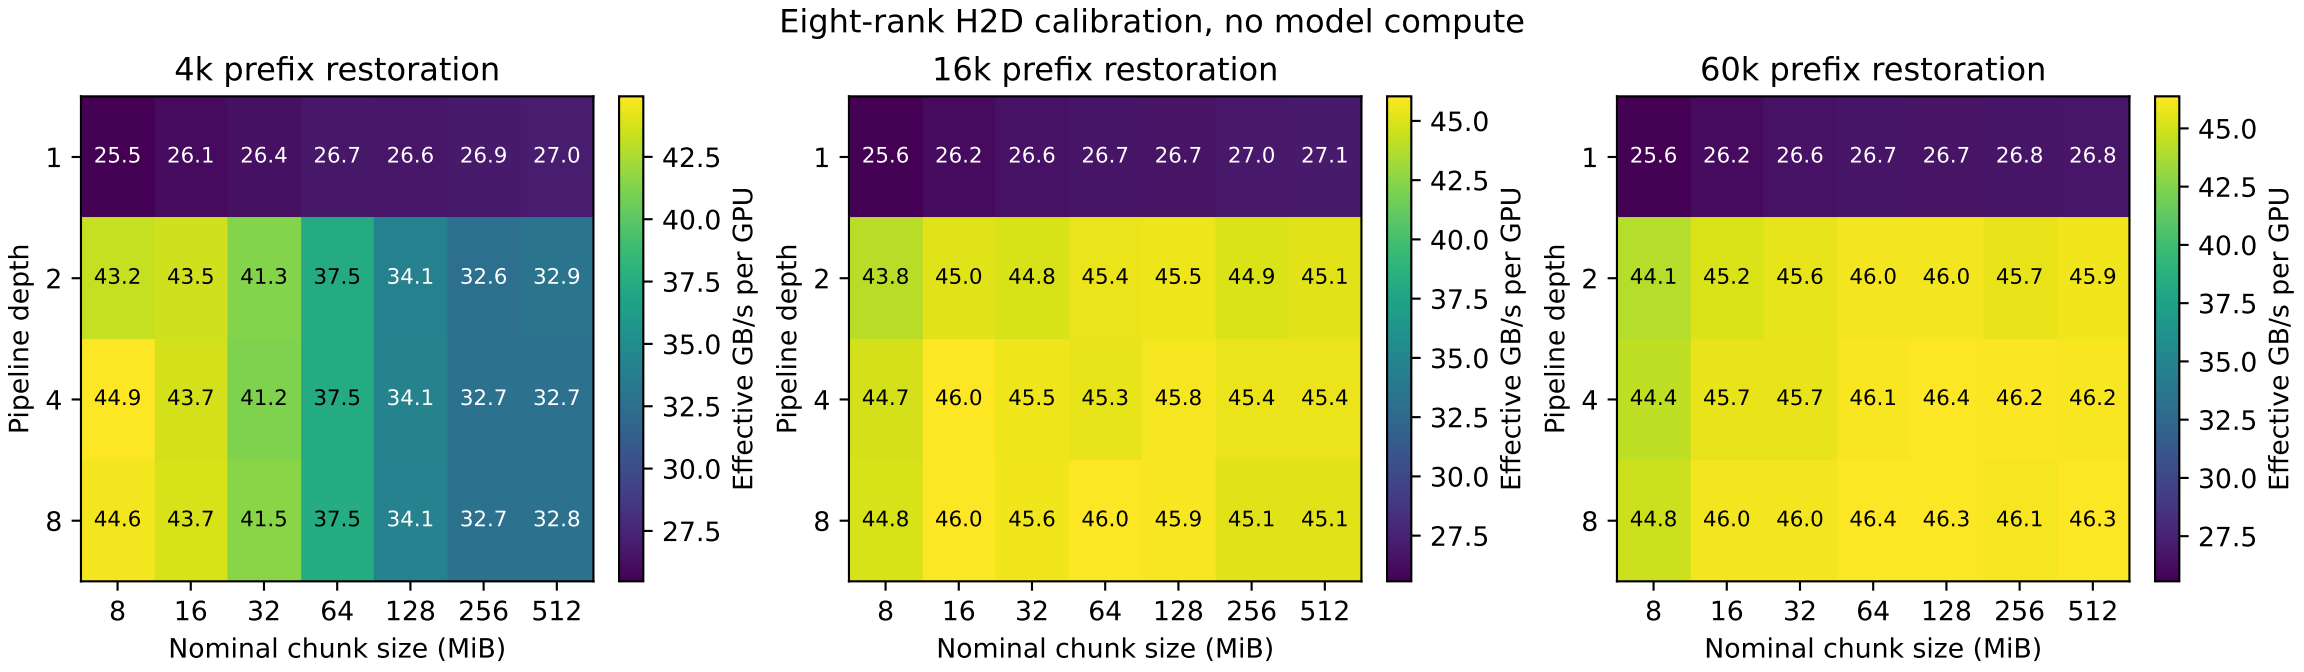
Eight-rank H2D calibration, no model compute
4k prefix restoration
16k prefix restoration
60k prefix restoration
Pipeline depth
Nominal chunk size (MiB)
Effective GB/s per GPU
1
2
4
8
8
16
32
64
128
256
512
25.5 26.1 26.4 26.7 26.6 26.9 27.0
43.2 43.5 41.3 37.5 34.1 32.6 32.9
44.9 43.7 41.2 37.5 34.1 32.7 32.7
44.6 43.7 41.5 37.5 34.1 32.7 32.8
25.6 26.2 26.6 26.7 26.7 27.0 27.1
43.8 45.0 44.8 45.4 45.5 44.9 45.1
44.7 46.0 45.5 45.3 45.8 45.4 45.4
44.8 46.0 45.6 46.0 45.9 45.1 45.1
25.6 26.2 26.6 26.7 26.7 26.8 26.8
44.1 45.2 45.6 46.0 46.0 45.7 45.9
44.4 45.7 45.7 46.1 46.4 46.2 46.2
44.8 46.0 46.0 46.4 46.3 46.1 46.3
27.5
30.0
32.5
35.0
37.5
40.0
42.5
45.0