KV readiness while compute runs · GLM 5.3 TP8 · eight B300 GPUs

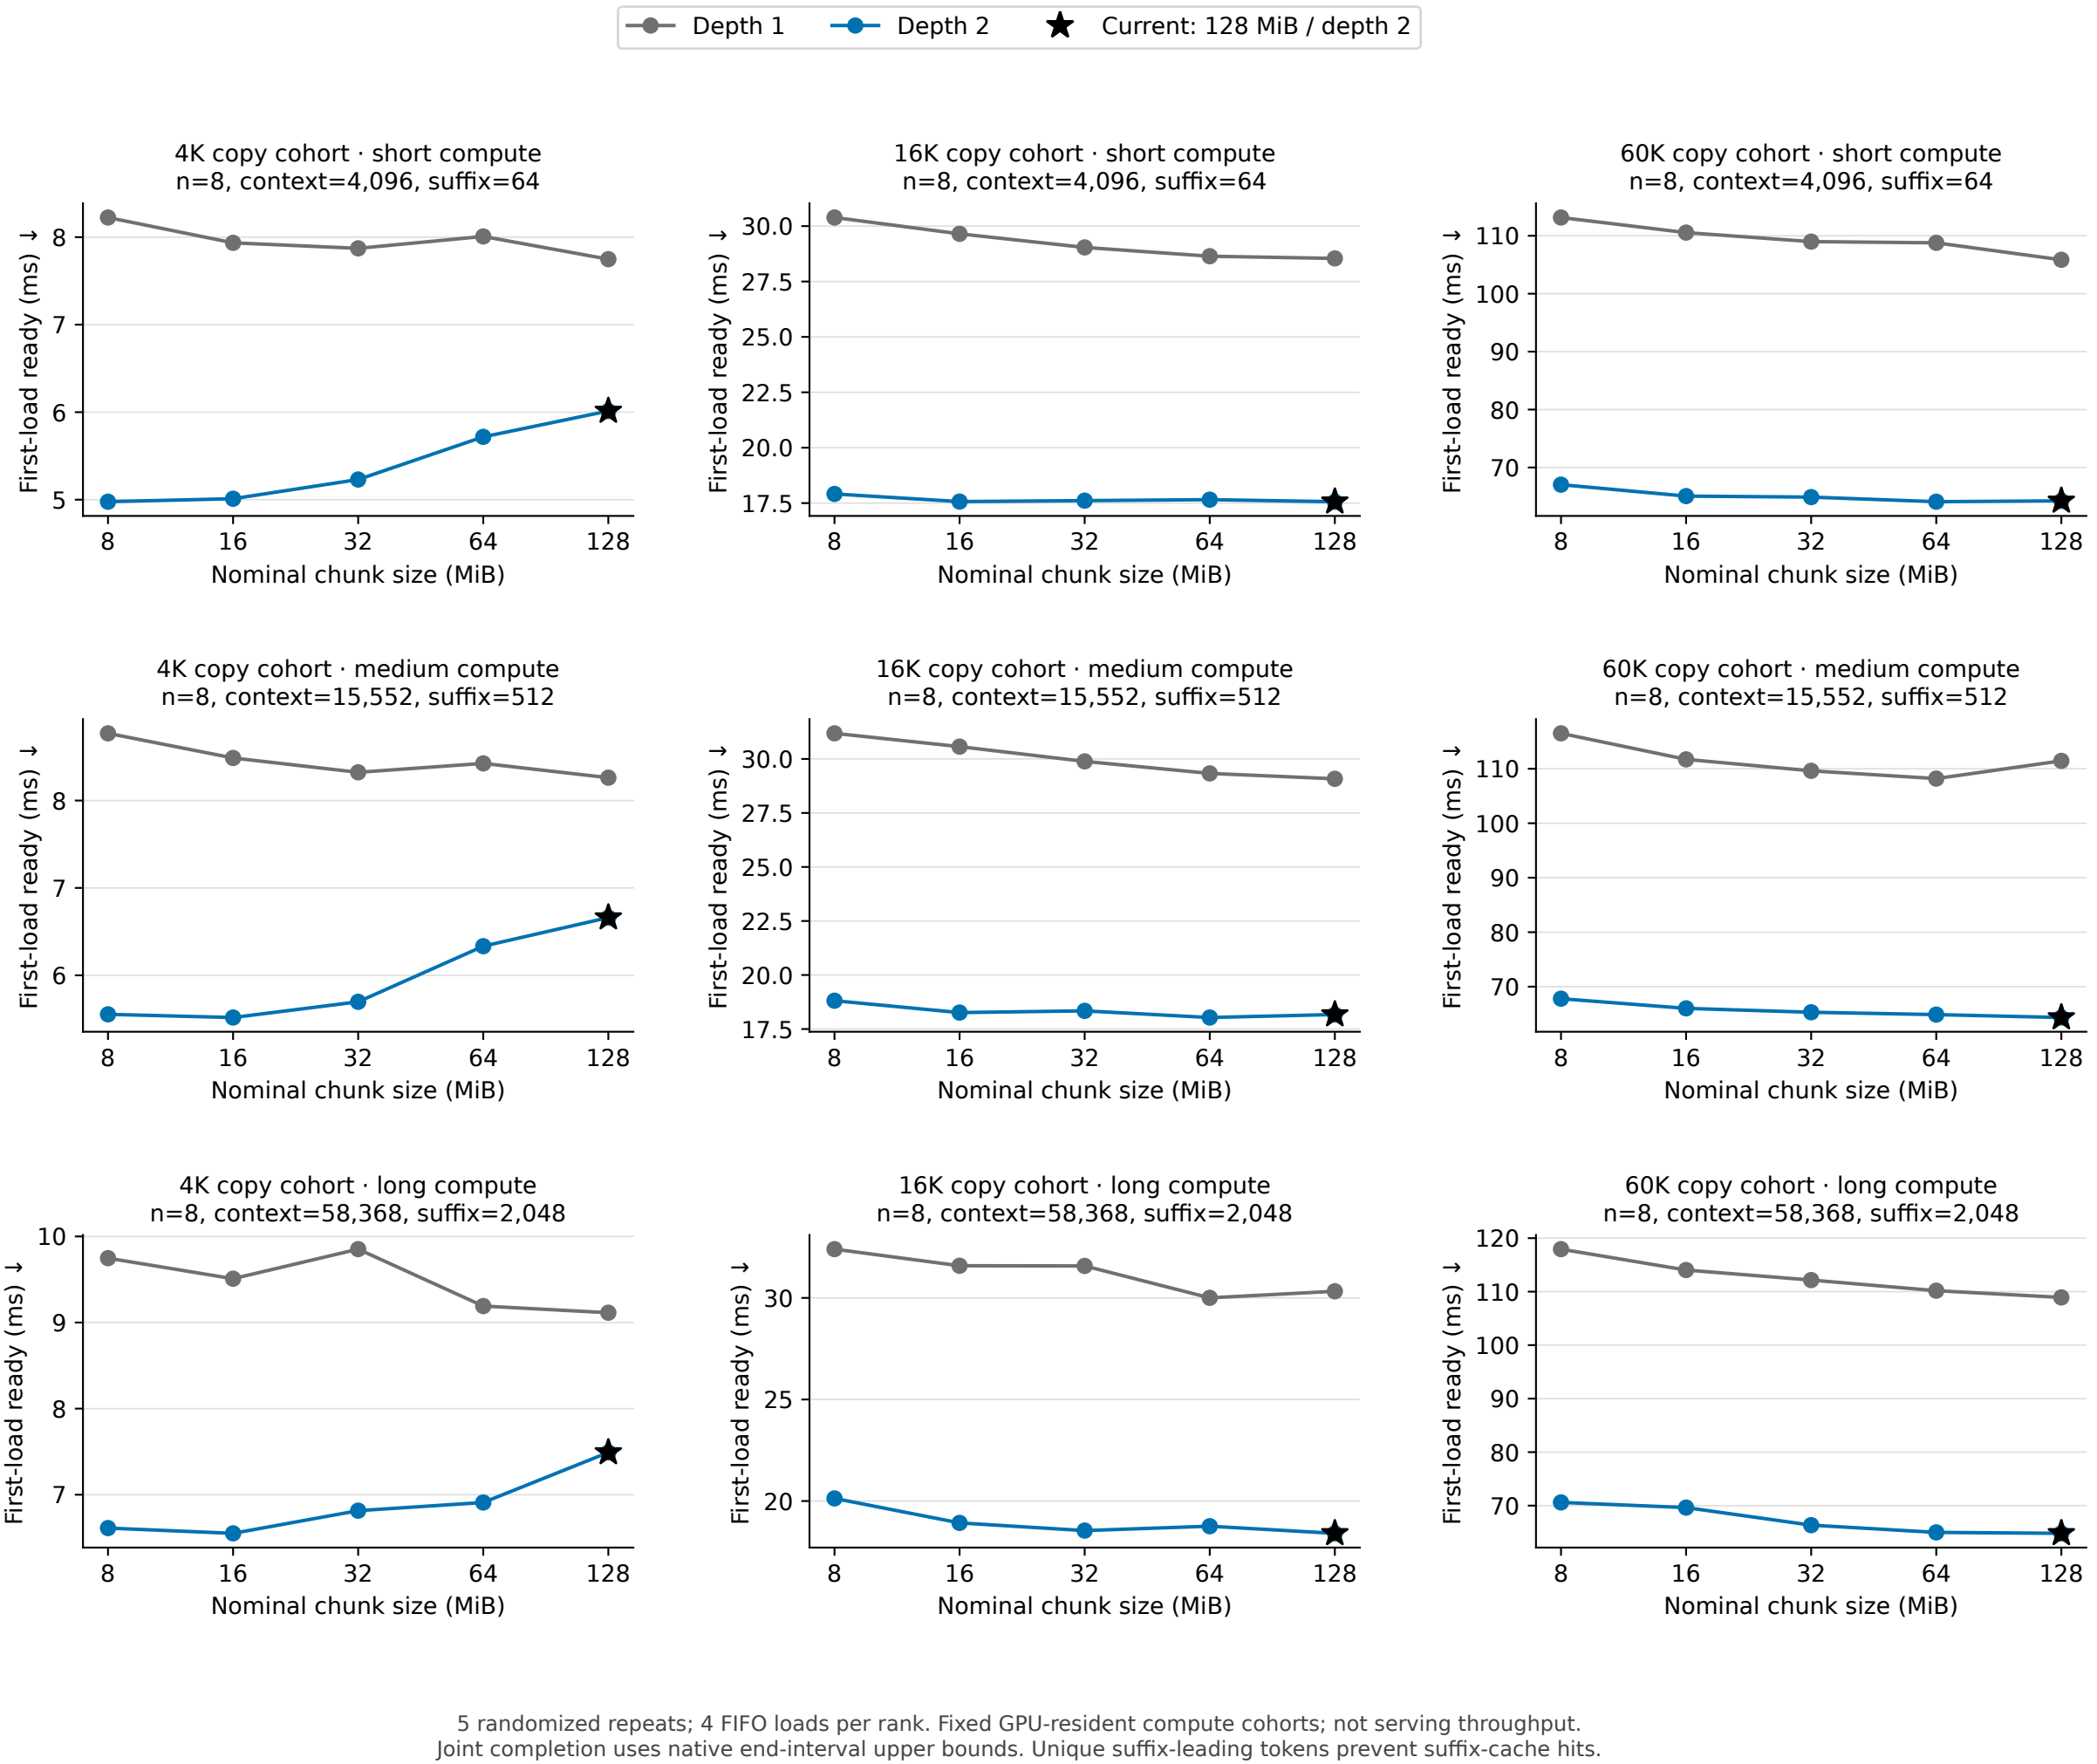


5 randomized repeats; 4 FIFO loads per rank. Fixed GPU-resident compute cohorts; not serving throughput.
Joint completion uses native end-interval upper bounds. Unique suffix-leading tokens prevent suffix-cache hits.