Host compute submission including copy enqueue · GLM 5.3 TP8 · eight B300 GPUs

Depth 1 · Depth 2 · Current: 128 MiB / depth 2

4K copy cohort · short compute
n=8, context=4,096, suffix=64

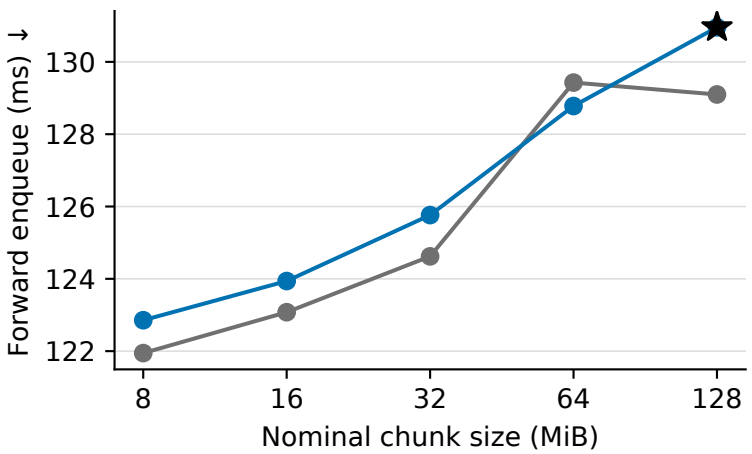

16K copy cohort · short compute
n=8, context=4,096, suffix=64

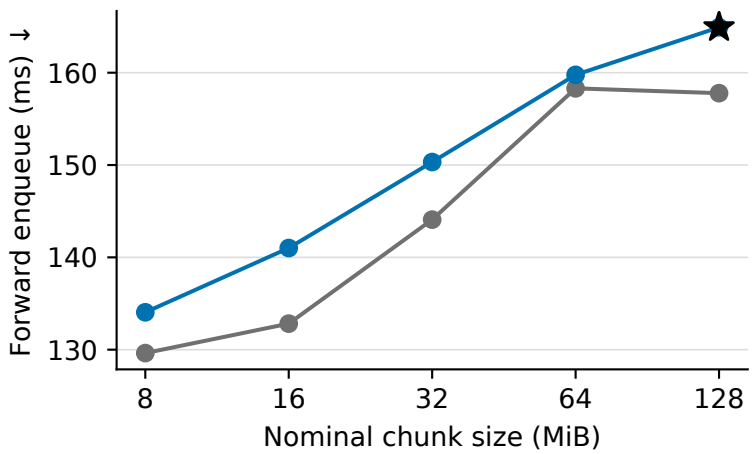

60K copy cohort · short compute
n=8, context=4,096, suffix=64

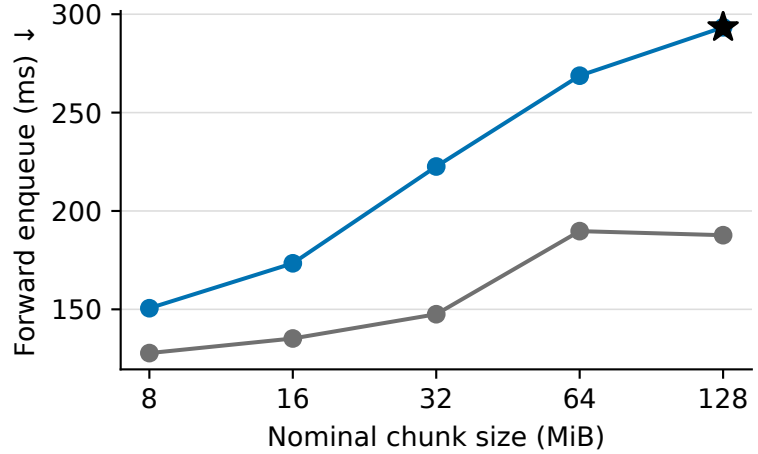

4K copy cohort · medium compute
n=8, context=15,552, suffix=512

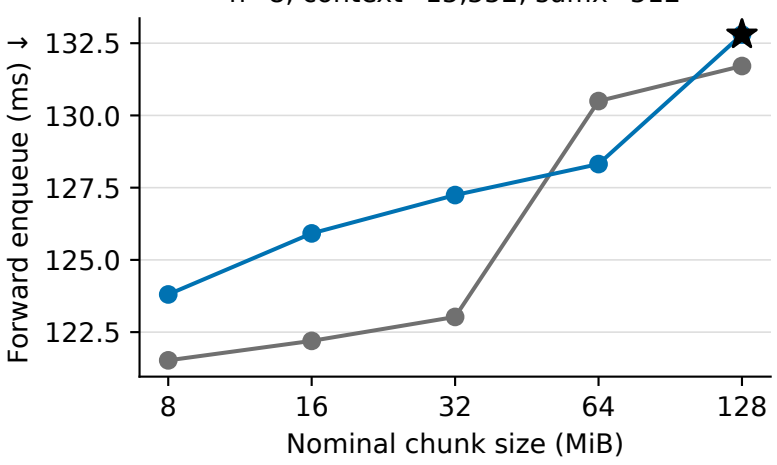

16K copy cohort · medium compute
n=8, context=15,552, suffix=512

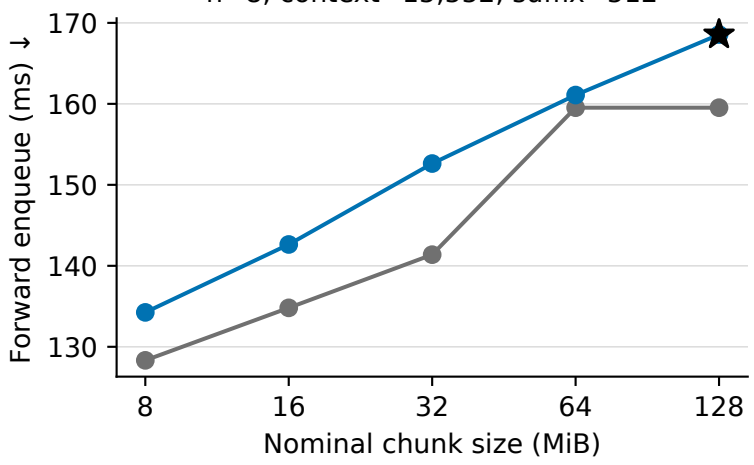

60K copy cohort · medium compute
n=8, context=15,552, suffix=512

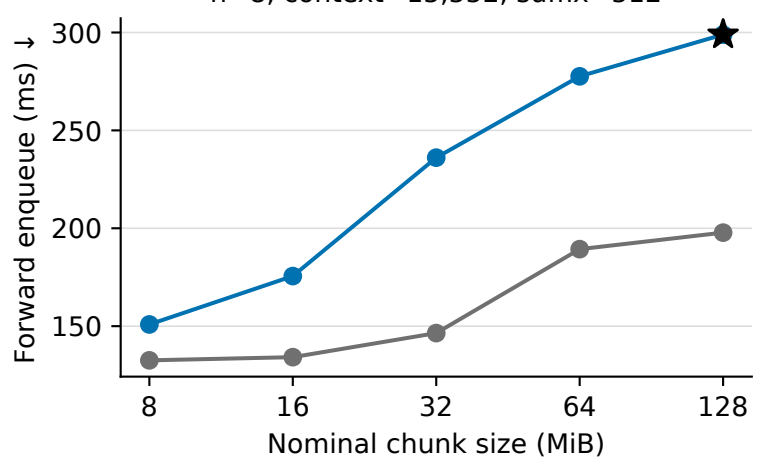

4K copy cohort · long compute
n=8, context=58,368, suffix=2,048

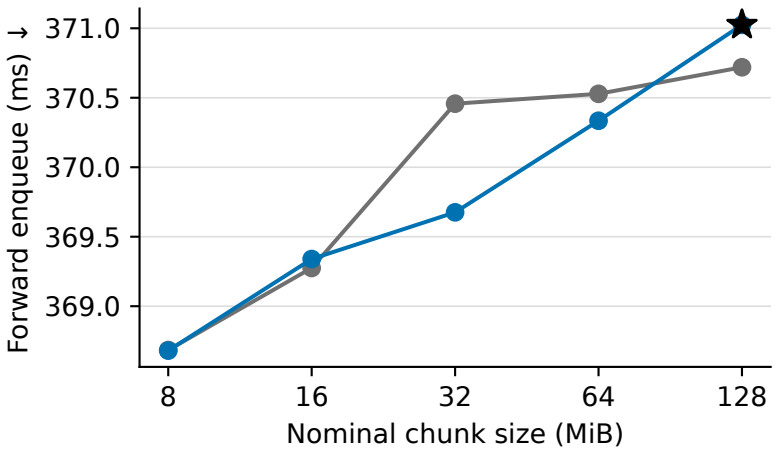

16K copy cohort · long compute
n=8, context=58,368, suffix=2,048

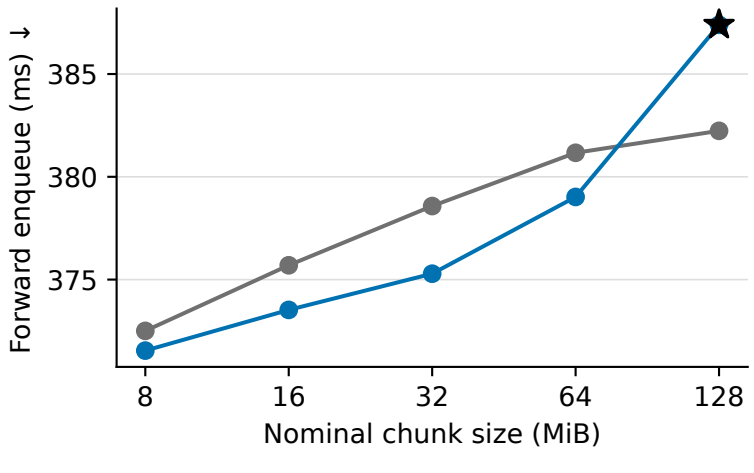

60K copy cohort · long compute
n=8, context=58,368, suffix=2,048

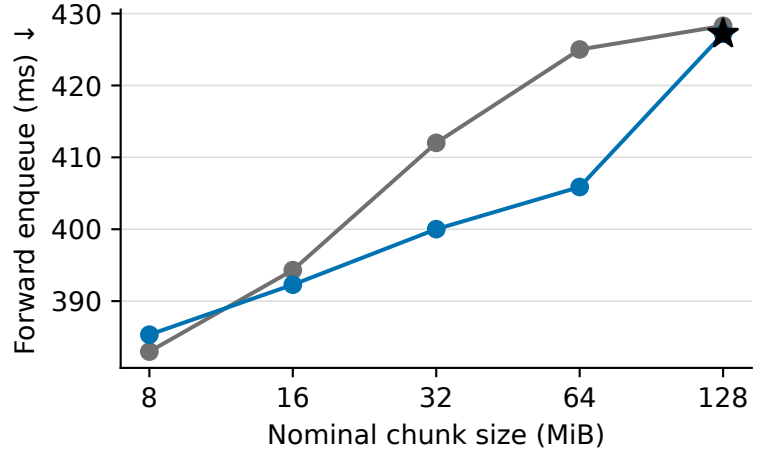

5 randomized repeats; 4 FIFO loads per rank. Fixed GPU-resident compute cohorts; not serving throughput.
Joint completion uses native end-interval upper bounds. Unique suffix-leading tokens prevent suffix-cache hits.