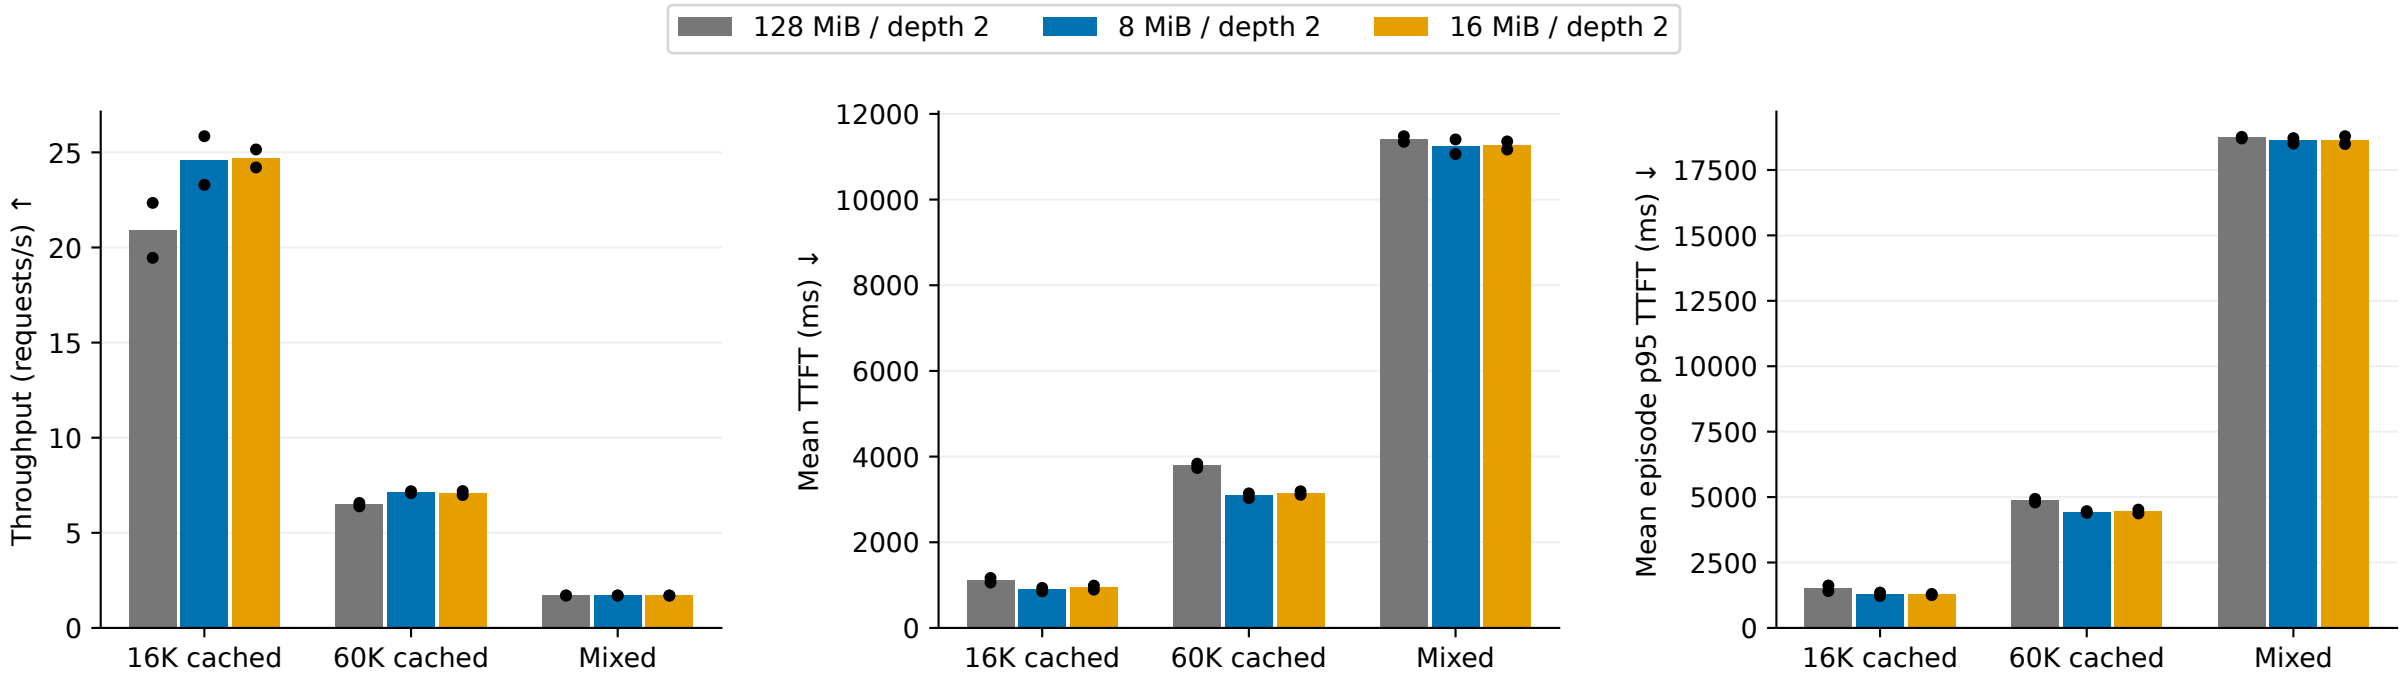
GLM 5.3 TP8 · unchanged latest scheduler / preparation policy
128 MiB / depth 2
8 MiB / depth 2
16 MiB / depth 2
Throughput (requests/s) ↑
0
5
10
15
20
25
16K cached
60K cached
Mixed
Mean TTFT (ms) ↓
0
2000
4000
6000
8000
10000
12000
16K cached
60K cached
Mixed
Mean episode p95 TTFT (ms) ↓
0
2500
5000
7500
10000
12500
15000
17500
16K cached
60K cached
Mixed
Two matched 32-request episodes per workload/config; concurrency 32; one output token. Bars: episode means; dots: episodes.
Separate allocations; small validation, not a continuous 300-request study or a statistical equivalence test.