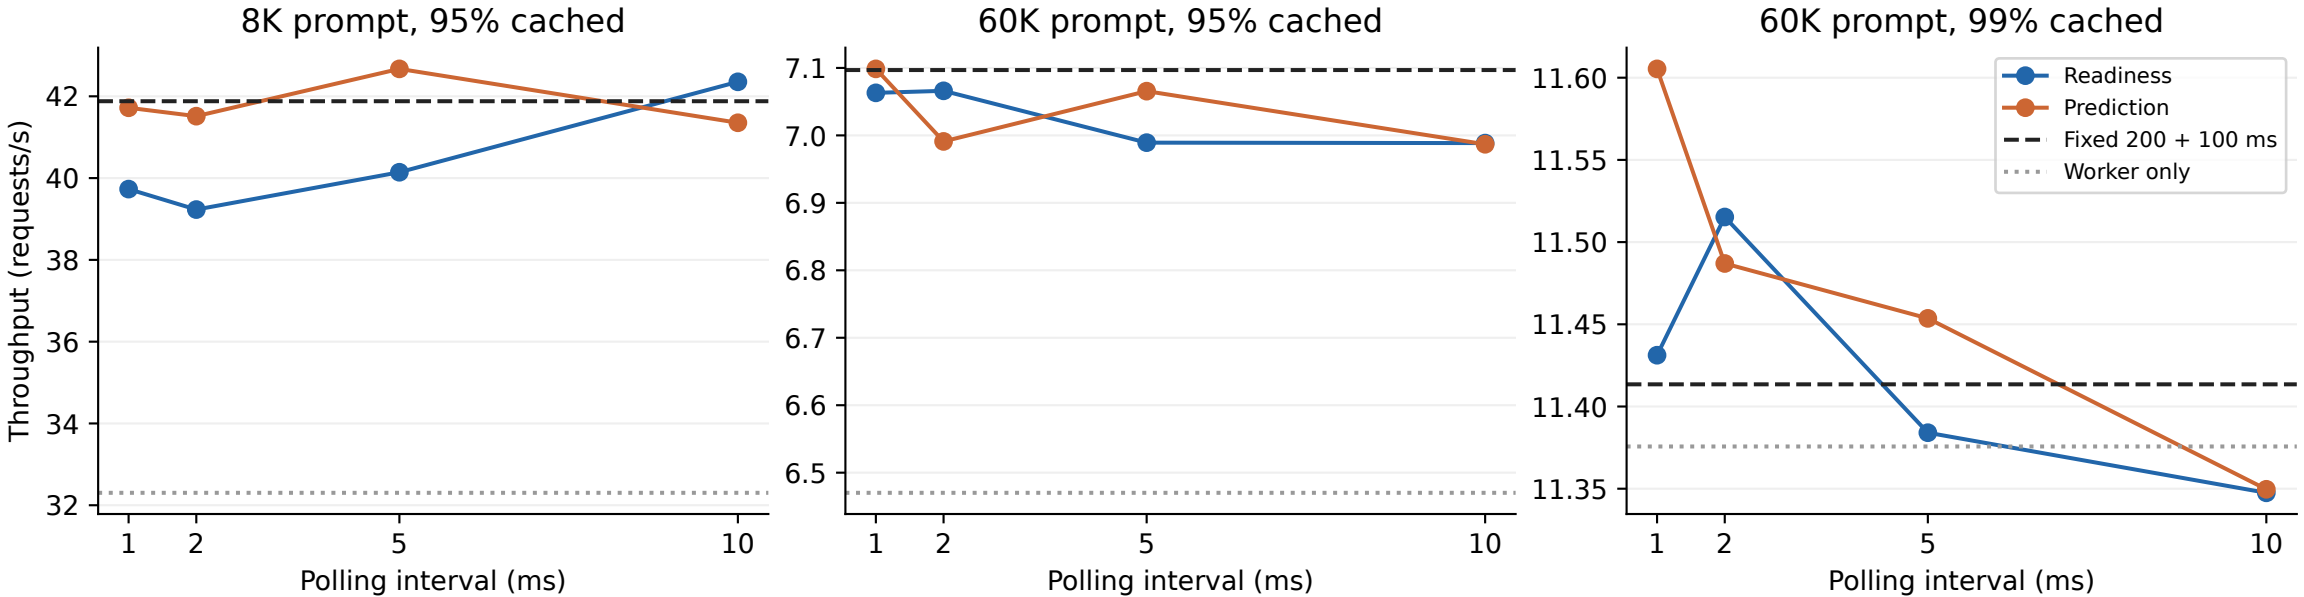
GLM 5.3 TP8 — screening cells, concurrency 32, one output token
8K prompt, 95% cached
60K prompt, 95% cached
60K prompt, 99% cached
Throughput (requests/s)
Polling interval (ms)
Readiness
Prediction
Fixed 200 + 100 ms
Worker only