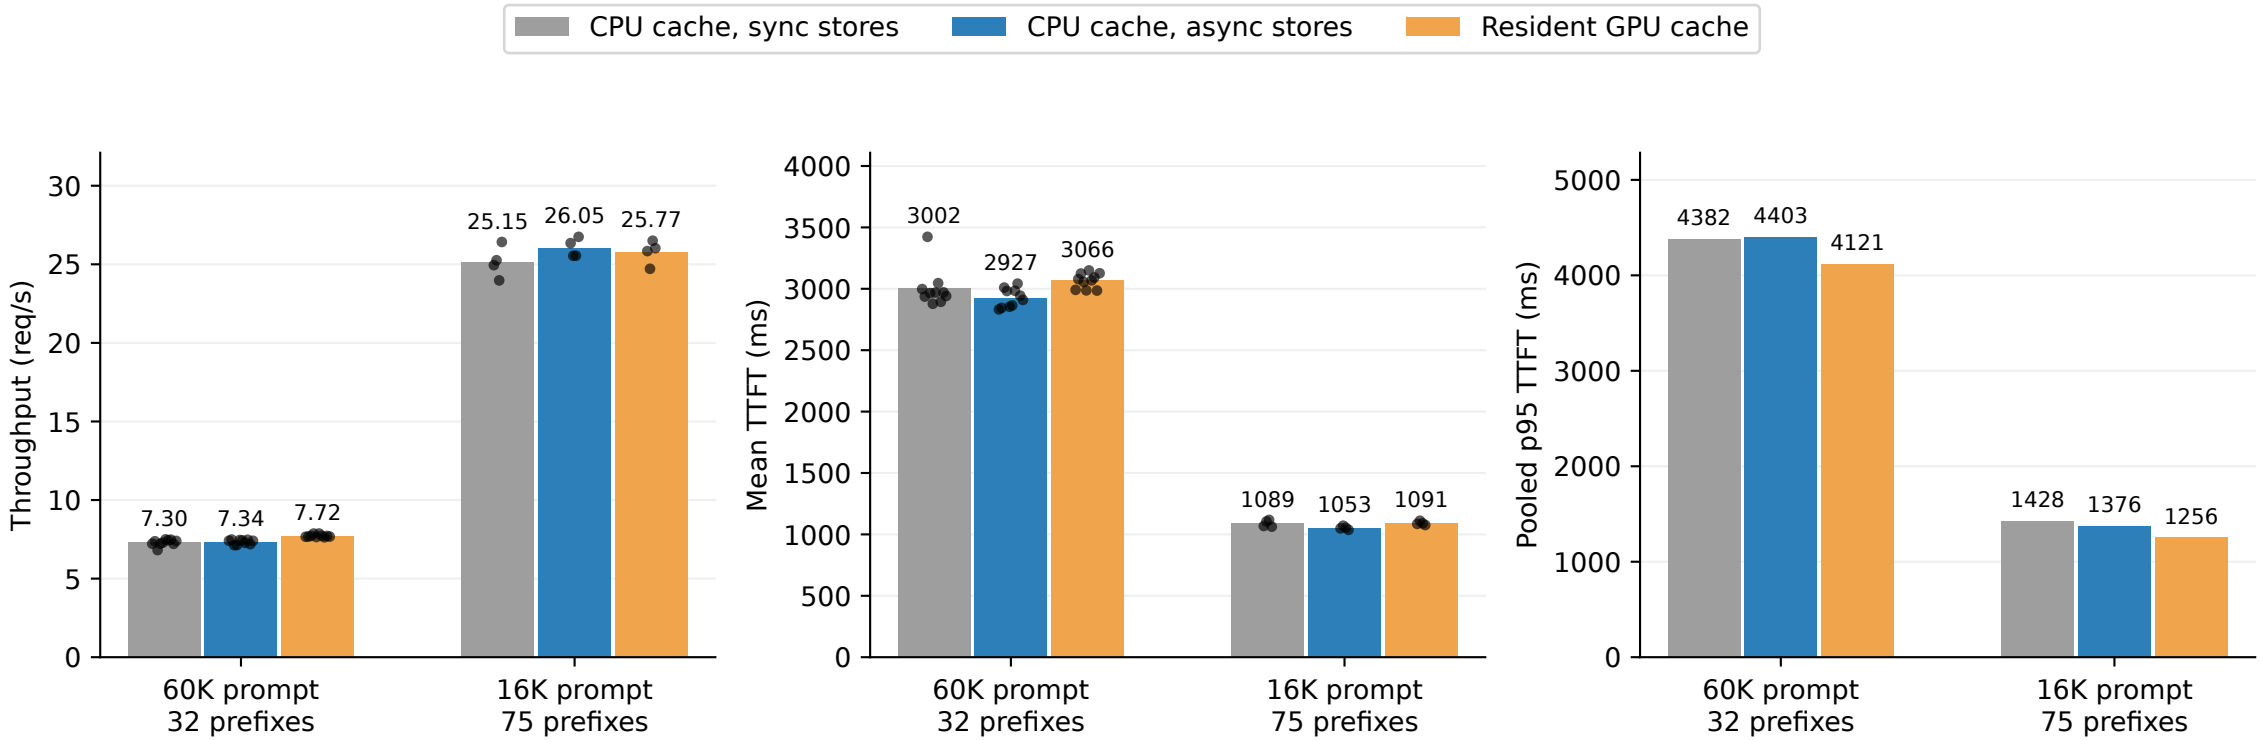
GLM 5.3 TP8: matched working sets that fit in GPU cache
CPU cache, sync stores
CPU cache, async stores
Resident GPU cache
Throughput (req/s)
7.30
7.34
7.72
25.15
26.05
25.77
Mean TTFT (ms)
3002
2927
3066
1089
1053
1091
Pooled p95 TTFT (ms)
4382
4403
4121
1428
1376
1256
60K prompt
32 prefixes
16K prompt
75 prefixes
Dots: episode repeats. Warm priming excluded. Fixed FP8 KV and cache capacity. GPU mode uses ordinary scheduling; a measured reference.