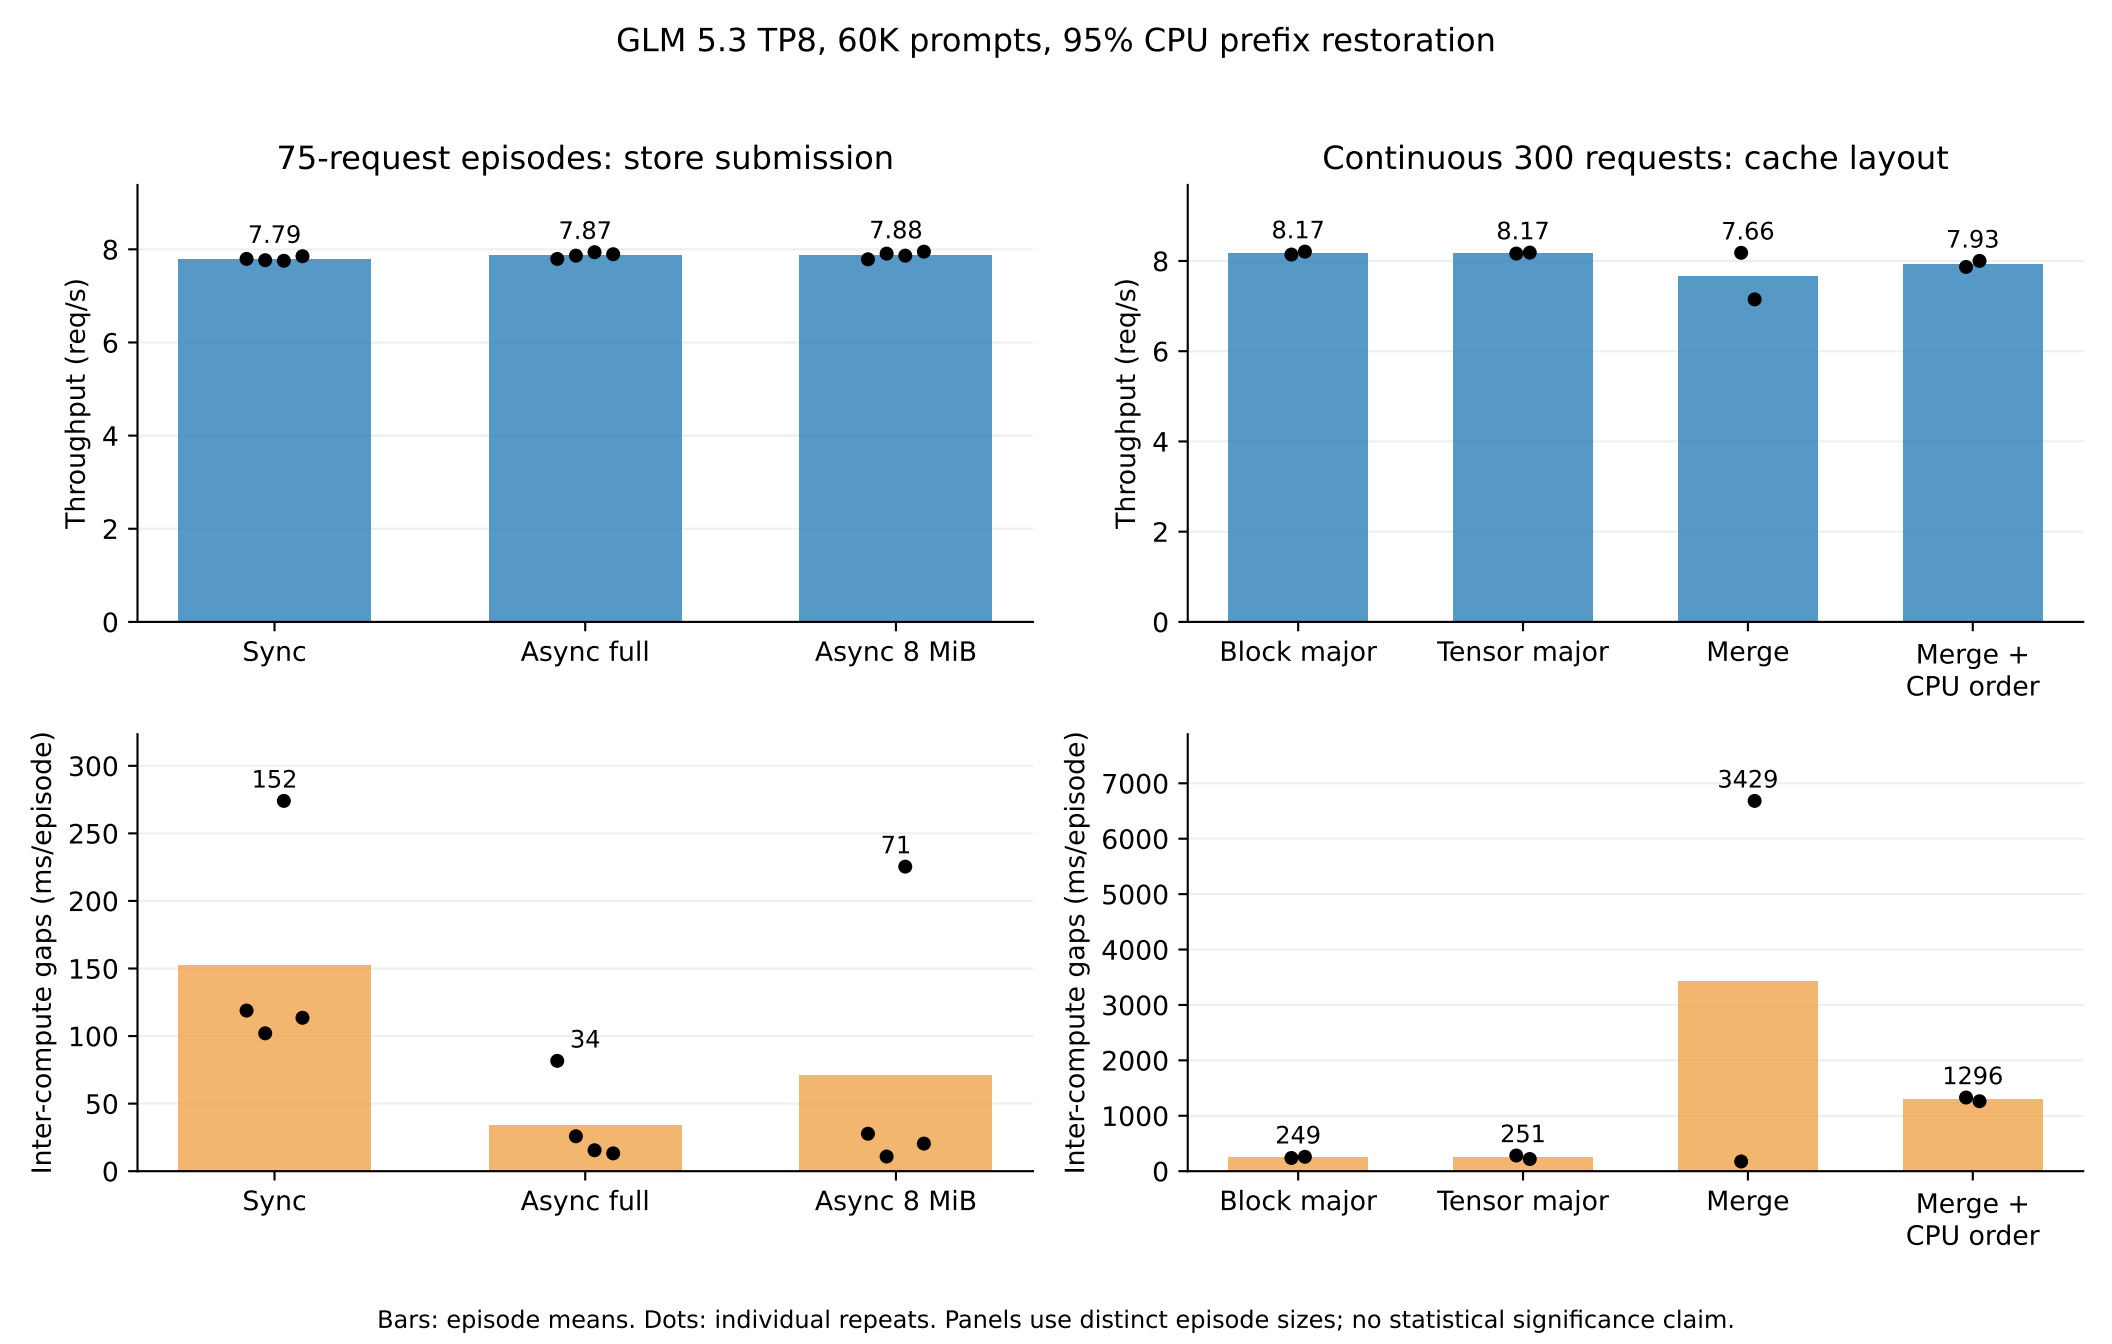
GLM 5.3 TP8, 60K prompts, 95% CPU prefix restoration
75-request episodes: store submission
Continuous 300 requests: cache layout
Throughput (req/s)
7.79
7.87
7.88
8.17
8.17
7.66
7.93
Sync
Async full
Async 8 MiB
Block major
Tensor major
Merge
Merge + CPU order
Inter-compute gaps (ms/episode)
152
34
71
249
251
3429
1296
Bars: episode means. Dots: individual repeats. Panels use distinct episode sizes; no statistical significance claim.